



AZ-204T00A

# Learning Path 09: Develop event-based solutions



# Module 2: Explore Azure Event Hubs

# Learning objectives

- Describe the benefits of using Event Hubs and how it captures streaming data.
- Explain how to process events.
- Perform common operations with the Event Hubs client library.

# Introduction

- Azure Event Hubs is a big data streaming platform and event ingestion service.
- It can receive and process millions of events per second.
- Data sent to an event hub can be transformed and stored by using any real-time analytics provider or batching/storage adapters.

# Discover Azure Event Hubs (1 of 2)

- Azure Event Hubs provides a unified streaming platform with time retention buffer, decoupling event producers from event consumers.
- It is a scalable event-processing service that ingests and processes large volumes of events and data, with low latency and high reliability.

| Feature                        | Description  |
|--------------------------------|--|
| Fully managed PaaS             | Event Hubs is a fully managed service with little configuration or management overhead   |
| Real-time and batch processing | Event Hubs uses a partitioned consumer model, enabling multiple applications to process the stream concurrently and letting you control the speed of processing. |
| Scalable                       | Scaling options, like Auto-inflate, scale the number of throughput units to meet your usage needs.   |
| Rich ecosystem                 | Event Hubs for Apache Kafka ecosystems enables Apache Kafka (1.0 and later) clients and applications to talk to Event Hubs                                       |


# Discover Azure Event Hubs (2 of 2)

## Key components

- An **Event Hub client** is the primary interface for developers interacting with the Event Hubs client library.
- An **Event Hub producer** is a type of client that serves as a source of telemetry data, diagnostics information, usage logs, or other log data, as part of an embedded device solution, a mobile device application, a game title running on a console or other device, some client or server based business solution, or a web site.
- An **Event Hub consumer** is a type of client which reads information from the Event Hub and allows processing of it. Processing may involve aggregation, complex computation and filtering.
- A **partition** is an ordered sequence of events that is held in an Event Hub. Partitions are a means of data organization associated with the parallelism required by event consumers.
- A **consumer group** is a view of an entire Event Hub. Consumer groups enable multiple consuming applications to each have a separate view of the event stream, and to read the stream independently at their own pace and from their own position.
- **Event receivers** - Any entity that reads event data from an event hub. All Event Hubs consumers connect via the AMQP 1.0 session.
- **Throughput units or processing units** - Pre-purchased units of capacity that control the throughput capacity of Event Hubs.

# Explore Event Hubs Capture (1 of 2)

- Azure Event Hubs enables you to automatically capture the streaming data in Event Hubs in an Azure Blob storage or Azure Data Lake Storage account of your choice.
- Event Hubs Capture enables you to process real-time and batch-based pipelines on the same stream.



# Explore Event Hubs Capture (2 of 2)

## Capture windowing

- Event Hubs Capture enables you to set up a window to control capturing.
- Each partition captures independently and writes a completed block blob at the time of capture, named for the time at which the capture interval was encountered.

## Scaling to throughput units

- Event Hubs traffic is controlled by throughput units.
- Event Hubs Capture copies data directly from the internal Event Hubs storage, bypassing throughput unit egress quotas and saving your egress for other processing readers.
- Event Hubs Capture runs automatically when you send your first event.

# Scale your processing application (1 of 2)

The key to scale for Event Hubs is the idea of *partitioned consumers*. In contrast to the competing consumers pattern, the partitioned consumer pattern enables high scale by removing the contention bottleneck and facilitating end to end parallelism.

## Example scenario

When designing the consumer in a distributed environment, the solution must handle the following requirements: scale, load balance, seamless resume on failures, and event consumption.

## Event processor or consumer client

- You don't need to build your own solution to meet these requirements.
- The Azure Event Hubs SDKs provide this functionality.
- For most production scenarios use the event processor client for reading and processing events.

# Scale your processing application (2 of 2)

## Partition ownership tracking

- An event processor instance typically owns and processes events from one or more partitions.
- Each event processor is given a unique identifier and claims ownership of partitions by adding or updating an entry in a checkpoint store.

## Checkpointing

- Checkpointing is a process by which an event processor marks or commits the position of the last successfully processed event within a partition.
- By specifying a lower offset from this checkpoint process, you can return to older data.

## Receive messages

- When you create an event processor, you specify the functions that will process events and errors.
- We recommend that you do things relatively fast. That is, do as little processing as possible.

## Thread safety and processor instances

By default, the function that processes the events is called sequentially for a given partition. As the event pump continues to run in the background of other threads, subsequent events from the same partition and calls to this function are queued in the background.

# Control access to events

## Azure built-in roles:

- Azure Event Hubs Data Owner: Use this role to give *complete access* to Event Hubs resources.
- Azure Event Hubs Data Sender: Use this role to give *send access* to Event Hubs resources.
- Azure Event Hubs Data Receiver: Use this role to give *receiving access* to Event Hubs resources.

## You can:

- Authorize access with managed identities
- Authorize access with Microsoft Identity Platform
- Authorize access to Event Hubs publishers with shared access signatures
- Authorize access to Event Hubs consumers with shared access signatures

# Perform common operations with the Event Hubs client library (1 of 3)

## Inspect an Event Hub

```
await using (var producer = new EventHubProducerClient(connectionString, eventHubName))
{
    string[] partitionIds = await producer.GetPartitionIdsAsync();
}
```

# Perform common operations with the Event Hubs client library (2 of 3)

## Publish events to an Event Hub

```
await using (var producer = new EventHubProducerClient(connectionString, eventHubName))
{
    using EventDataBatch eventBatch = await producer.CreateBatchAsync();
    eventBatch.TryAdd(new EventData(new BinaryData("First")));
    eventBatch.TryAdd(new EventData(new BinaryData("Second")));

    await producer.SendAsync(eventBatch);
}
```

# Perform common operations with the Event Hubs client library (3 of 3)

## Read events from an Event Hub

```
string consumerGroup = EventHubConsumerClient.DefaultConsumerGroupName;

await using (var consumer = new EventHubConsumerClient(consumerGroup, connectionString, eventHubName))
{
    using var cancellationSource = new CancellationTokenSource();
    cancellationSource.CancelAfter(TimeSpan.FromSeconds(45));

    await foreach (PartitionEvent receivedEvent in consumer.ReadEventsAsync(cancellationSource.Token))
    {
        // At this point, the loop will wait for events to be available in the Event Hub. When an event
        // is available, the loop will iterate with the event that was received. Because we did not
        // specify a maximum wait time, the loop will wait forever unless cancellation is requested using
        // the cancellation token.
    }
}
```

# Exercise – Send and retrieve events from Azure Event Hubs



Create Azure Event Hubs resources and build a .NET console app to send and receive events using the `Azure.Messaging.EventHubs` SDK.

Use the Lab [AZ-104T00-A-CEP-OP] M01 Manage Azure Active Directory Identities

- <https://microsoftlearning.github.io/mslearn-azure-developer/instructions/azure-events-messages/02-event-hubs-send-receive.html>
- <https://learn.microsoft.com/en-us/azure/connectors/connectors-create-api-azure-event-hubs>

# Module 2 knowledge check

1

Which Event Hubs concept represents an ordered sequence of events that is held in an Event Hubs?

2

Which process represents when an event processor marks or commits the position of the last successfully processed event within a partition?

# Module 2 summary

In this module, you learned how to:

- Describe the benefits of using Event Hubs and how it captures streaming data.
- Explain how to process events.
- Perform common operations with the Event Hubs client library.

# Group discussion

# Group discussion questions

- Can you describe the differences between the capabilities of Azure Event Grid and Azure Event Hubs? When would you use one over the other?
- Contoso Inc is using Event Grid in their solution, but many events are being lost during delivery. What can they do to improve delivery and retain the events being lost?
- Can you list the three Azure built-in roles for Event Hubs and what permissions they grant?

# End of presentation